

Application Design Using Java

Lecture 20

Database Systems

- *DBMS (Database Management System)* is a software tool for storing and managing large amounts of data
- *Database* is a collection of data organized for a specific application, often stored in a DBMS
- *Database application* is a software product that uses DBMSs to store one or more databases to accomplish a specific purpose



What *Is* a Relational Database Management System?

Database Management System = DBMS

Relational DBMS = RDBMS

- A collection of files that store the data
- A big C program written by someone else that accesses and updates those files for you

Where are RDBMS used?

- Backend for traditional “database” applications
- Backend for large Websites
- Backend for Web services

Example of a Traditional Database Application

Suppose we are building a system
to store the information about:

- students
- courses
- professors
- who takes what, who teaches what

Can we do it without a DBMS ?

Sure we can! Start by storing the data in files:

students.txt

courses.txt

professors.txt



Now write C or Java programs to implement specific tasks

Doing it without a DBMS...

- Enroll “Mary Johnson” in “CSE444”:

Write a C program to do the following:

Read ‘students.txt’

Read ‘courses.txt’

Find&update the record “Mary Johnson”

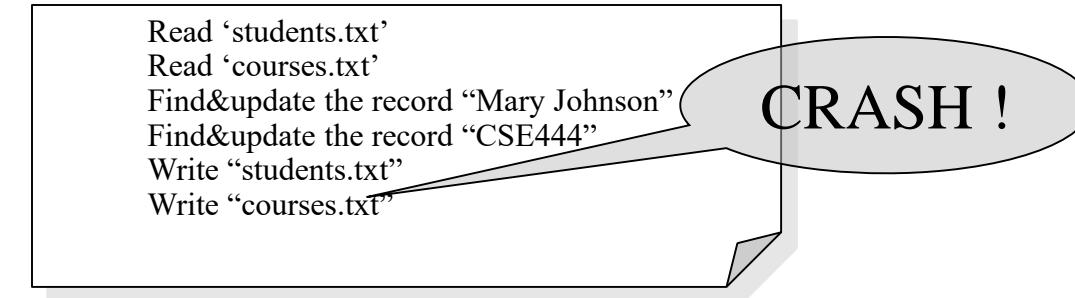
Find&update the record “CSE444”

Write “students.txt”

Write “courses.txt”

Problems without an DBMS...

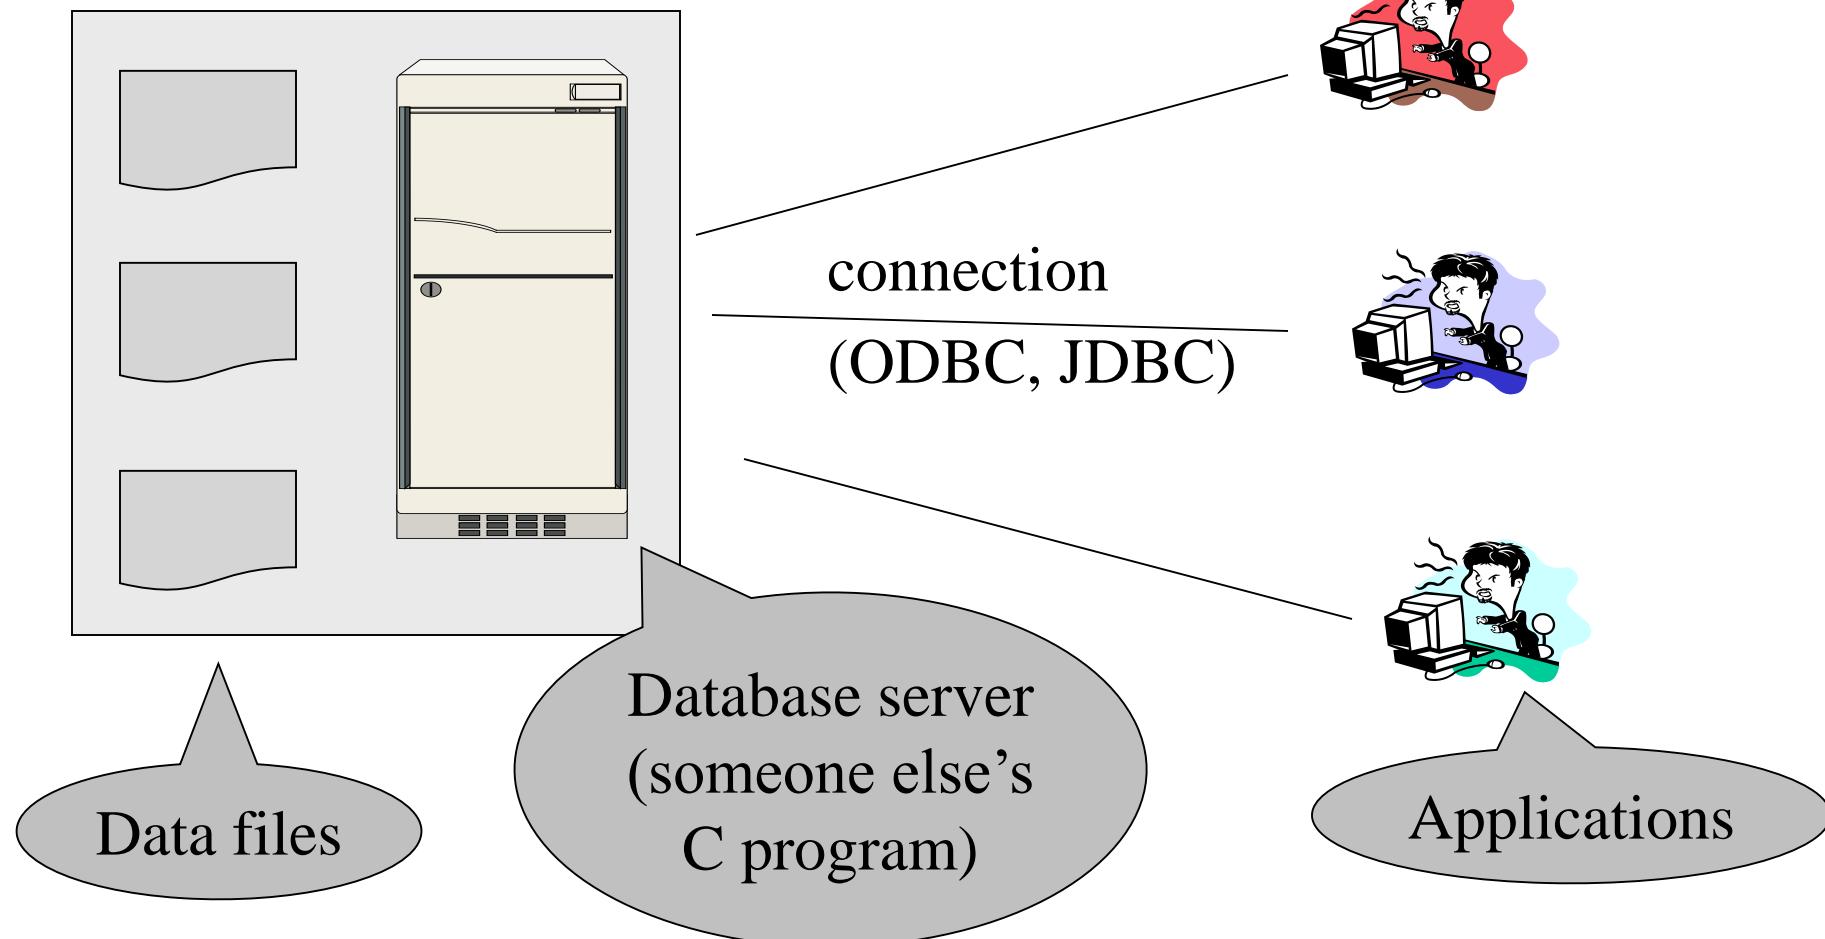
- System crashes:



- What is the problem ?
- Large data sets (say, 50GB)
 - What is the problem ?
- Simultaneous access by many users
 - Need locks: we know them from OS, but now data on disk; and is there any fun to re-implement them ?

Enters a DMBS

“Two tier database system”



Functionality of a DBMS

The programmer sees SQL, which has two components:

- Data Definition Language - DDL
- Data Manipulation Language - DML
 - query language

Behind the scenes the DBMS has:

- Query optimizer
- Query engine
- Storage management
- Transaction Management (concurrency, recovery)

Functionality of a DBMS

Two things to remember:

- Client-server architecture
 - Slow, cumbersome connection
 - But good for the data
- It is just someone else's C program
 - In the beginning we may be impressed by its speed
 - But later we discover that it can be frustratingly slow
 - We can do any particular task faster outside the DBMS
 - But the DBMS is *general* and *convenient*

How the Programmer Sees the DBMS

- Start with DDL to *create tables*:

```
CREATE TABLE Students (
    Name CHAR(30)
    SSN CHAR(9) PRIMARY KEY NOT NULL,
    Category CHAR(20)
) ...
```

- Continue with DML to *populate tables*:

```
INSERT INTO Students
VALUES('Charles', '123456789', 'undergraduate')
. . .
```

How the Programmer Sees the DBMS

- Tables:

Students:

SSN	Name	Category
123-45-6789	Charles	undergrad
234-56-7890	Dan	grad

Takes:

SSN	CID
123-45-6789	CSE444
123-45-6789	CSE444
234-56-7890	CSE142
	...

Courses:

CID	Name	Quarter
CSE444	Databases	fall
CSE541	Operating systems	winter

- Still implemented as files, but behind the scenes can be quite complex

“*data independence*” = separate *logical view* from *physical implementation*

Transactions

- Enroll “Mary Johnson” in “CSE444”:

```
BEGIN TRANSACTION;

INSERT INTO Takes
    SELECT Students.SSN, Courses.CID
    FROM Students, Courses
    WHERE Students.name = 'Mary Johnson' and
          Courses.name = 'CSE444'

    -- More updates here.....

    IF everything-went-OK
        THEN COMMIT;
    ELSE ROLLBACK
```

If system crashes, the transaction is still either committed or aborted

Transactions

- A *transaction* = sequence of statements that either all succeed, or all fail
- Transactions have the ACID properties:
 - A = atomicity
 - C = consistency
 - I = independence
 - D = durability

Queries

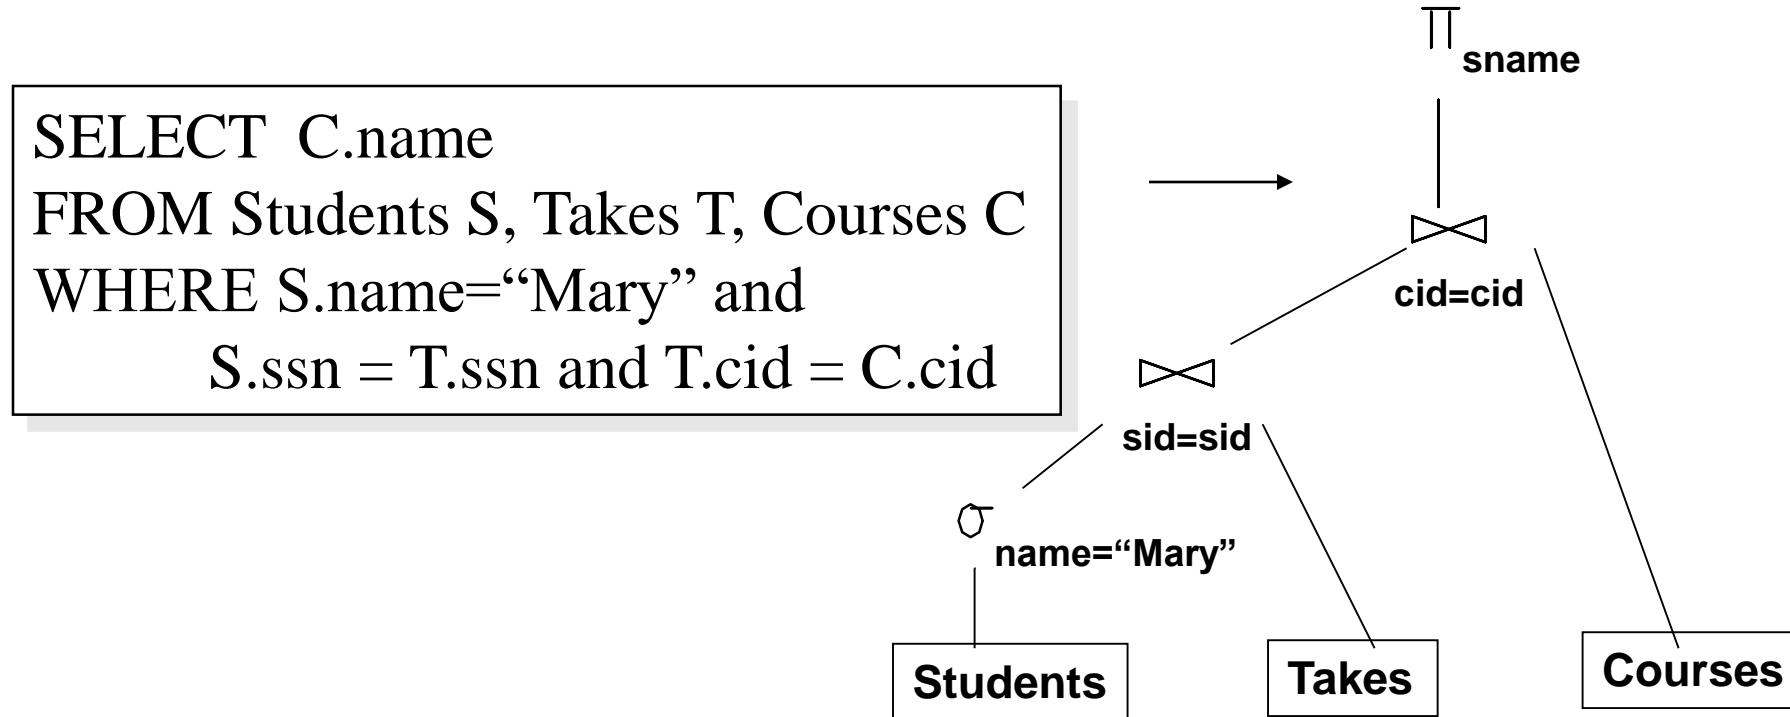
- Find all courses that “Mary” takes

```
SELECT C.name  
FROM Students S, Takes T, Courses C  
WHERE S.name = "Mary" and  
S.ssn = T.ssn and T.cid = C.cid
```

- What happens behind the scene ?
 - Query processor figures out how to answer the query efficiently.

Queries, behind the scene

Declarative SQL query → *Imperative query execution plan:*



The **optimizer** chooses the best execution plan for a query

Database Systems

- The big commercial database vendors:
 - Oracle
 - IBM (with DB2) bought Informix recently
 - Microsoft (SQL Server)
 - Sybase
- Some free database systems (Unix) :
 - Postgres
 - MySQL
 - Predator
- We will use MySQL

New Trends in Databases

- Object-relational databases
- Main memory database systems
- XML XML XML !
 - Relational databases with XML support
 - Middleware between XML and relational databases
 - Native XML database systems
 - Lots of research here at UW on XML and databases
- Data integration
- Peer to peer, stream data management – still research

SQL Introduction

- Standard language for querying and manipulating data
- Structured Query Language
 - Many standards out there:
 - ANSI SQL
 - SQL92 (a.k.a. SQL2)
 - SQL99 (a.k.a. SQL3)
 - Vendors support various subsets of these
 - What we discuss is common to all of them

SQL

- Data Definition Language (DDL)
 - Create/alter/delete tables and their attributes
- Data Manipulation Language (DML)
 - Query one or more tables
 - Insert/delete/modify tuples in tables

Data in SQL

1. Atomic types, a.k.a. data types
2. Tables built from atomic types

Unlike XML, no nested tables, only flat tables are allowed!

- We will see later how to decompose complex structures into multiple flat tables

Data Types in SQL

- Characters:
 - CHAR(20) -- fixed length
 - VARCHAR(40) -- variable length
- Numbers:
 - BIGINT, INT, SMALLINT, TINYINT
 - REAL, FLOAT -- differ in precision
 - MONEY
- Times and dates:
 - DATE
 - DATETIME -- SQL Server
- Others... All are simple

Tables in SQL

- Table name
- Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

- Attribute names

- Tuples or rows

Tables Explained

- A tuple = a record
 - Restriction: all attributes are of atomic type
- A table = a set of tuples
 - Like a list...
 - ...but it is unordered: no **first()**, no **next()**, no **last()**.

Tables Explained

- The *schema* of a table is the table name and its attributes:

Product(PName, Price, Category, Manufacturer)

- A *key* is an attribute whose values are unique; we underline a key

Product(PName, Price, Category, Manufacturer)

SQL Query

- Basic form: (plus many many more bells and whistles)

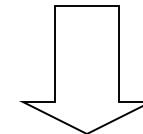
- **SELECT** attributes
- **FROM** relations (possibly multiple)
- **WHERE** conditions (selections)

Simple SQL Query

•Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
•SELECT *
  FROM Product
 WHERE Category='Gadgets'
```



•“selection”

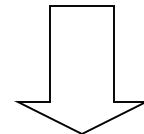
PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks

Simple SQL Query

•Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
•SELECT PName, Price, Manufacturer  
FROM Product  
WHERE Price > 100
```

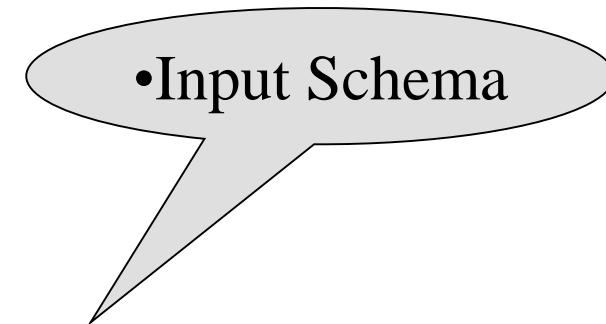


- “selection” and
- “projection”

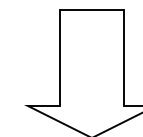
PName	Price	Manufacturer
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

A Notation for SQL Queries

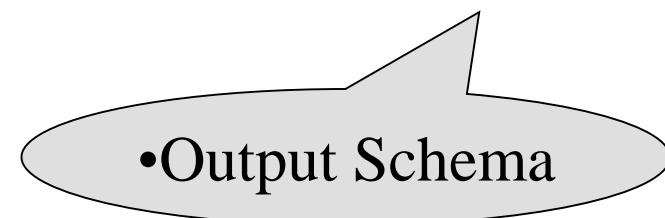
```
•SELECT PName, Price, Manufacturer  
FROM Product  
WHERE Price > 100
```



•Product(PName, Price, Category, Manufacturer)



•Answer(PName, Price, Manufacturer)



•Output Schema

Selections

What goes in the **WHERE** clause:

- $x = y$, $x < y$, $x \leq y$, etc
 - For numbers, they have the usual meanings
 - For CHAR and VARCHAR: lexicographic ordering
 - Expected conversion between CHAR and VARCHAR
 - For dates and times, what you expect...
- Pattern matching on strings...

The LIKE operator

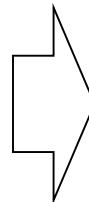
- $s \text{ LIKE } p$: pattern matching on strings
- p may contain two special symbols:
 - $\%$ = any sequence of characters
 - $_$ = any single character

Product(PName, Price, Category, Manufacturer)
Find all products whose name mentions 'gizmo':

```
•SELECT *
FROM   Products
WHERE  PName LIKE '%gizmo%'
```

Eliminating Duplicates

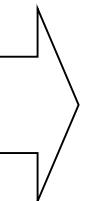
```
•SELECT DISTINCT Category  
•FROM Product
```



Category
Gadgets
Photography
Household

•Compare to:

```
•SELECT Category  
•FROM Product
```



Category
Gadgets
Gadgets
Photography
Household

Ordering the Results

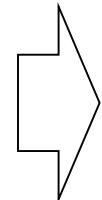
```
•SELECT PName, Price, Manufacturer  
FROM Product  
WHERE Category='gizmo' AND Price > 50  
ORDER BY Price, PName
```

- Ordering is ascending, unless you specify the DESC keyword.
- Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering the Results

```
•SELECT Category  
•FROM Product  
•ORDER BY PName
```

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi



• ?

Ordering the Results

```
•SELECT DISTINCT Category  
FROM Product  
ORDER BY Category
```

Category
Gadgets
Household
Photography

•Compare to:

```
•SELECT DISTINCT Category  
FROM Product  
ORDER BY PName
```



Joins in SQL

- Connect two or more tables:

•Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

•Company

•What is
the connection
between
them ?

CName	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

Joins

- Product (PName, Price, Category, Manufacturer)
- Company (CName, StockPrice, Country)

- Find all products under \$200 manufactured in Japan; return their names and prices.

```
•SELECT PName, Price  
FROM Product, Company  
WHERE Manufacturer=CName AND Country='Japan'  
      AND price <= 200
```

• Join
between Product
and Company

Joins in SQL

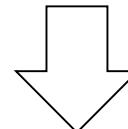
•Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

•Company

Cname	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
•SELECT PName, Price  
FROM Product, Company  
WHERE Manufacturer=CName AND country='Japan'  
AND price <= 200
```



PName	Price
SingleTouch	\$149.99

Joins

- Product (PName, Price, Category, Manufacturer)
- Company (CName, StockPrice, Country)
-
- Find all countries that manufacture some product in the ‘Gadgets’ category.

```
•SELECT Country  
FROM Product, Company  
WHERE Manufacturer=CName AND Category='Gadgets'
```

Joins in SQL

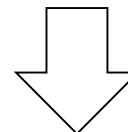
•Product

Name	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

•Company

Cname	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
•SELECT Country  
FROM Product, Company  
WHERE Manufacturer=CName AND Category='Gadgets'
```



Country
??
??

Joins

- Product (PName, Price, Category, Manufacturer)
 - Purchase (Buyer, Seller, Store, Product)
 - Person(Persname, PhoneNumber, City)
-
- Find names of people living in Seattle that bought some product in the ‘Gadgets’ category, and the names of the stores they bought such product from

```
•SELECT DISTINCT Persname, Store  
FROM Person, Purchase, Product  
WHERE Persname=Buyer AND Product = Pname AND  
City='Seattle' AND Category='Gadgets'
```

When are two tables related?

- You guess they are
- I tell you so
- Foreign keys are a method for schema designers to tell you so
 - A foreign key states that a column is a reference to the key of another table
ex: **Product.Manufacturer** is foreign key of **Company**
 - Gives information and enforces constraint

Disambiguating Attributes

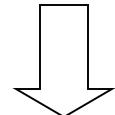
- Sometimes two relations have the same attribute name:

Person(PName, Address, Worksfor)

Company(CName, Address)

```
•SELECT DISTINCT PName, Address  
FROM Person, Company  
WHERE Worksfor = CName
```

•Which
address ?



```
•SELECT DISTINCT Person.PName, Company.Address  
FROM Person, Company  
WHERE Person.Worksfor = Company.CName
```

Tuple Variables

- Product (PName, Price, Category, Manufacturer)
- Purchase (Buyer, Seller, Store, Product)
- Person(Persname, PhoneNumber, City)
- Find all stores that sold at least one product that the store ‘BestBuy’ also sold:

```
•SELECT DISTINCT x.store  
•FROM Purchase AS x, Purchase AS y  
•WHERE x.product = y.product AND y.store = 'BestBuy'
```

- Answer (store)

Tuple Variables

- General rule:

tuple variables introduced automatically by the system:

- Product (name, price, category, manufacturer)

```
• SELECT Name  
FROM Product  
WHERE Price > 100
```

- Becomes:

```
• SELECT Product.Name  
FROM Product AS Product  
WHERE Product.Price > 100
```

- Doesn't work when Product occurs more than once:
- In that case the user needs to define variables explicitly.

Meaning (Semantics) of SQL Queries

SELECT a₁, a₂, ..., a_k
FROM R₁ AS x₁, R₂ AS x₂, ..., R_n AS x_n

WHERE Conditions

1. Nested loops:

```
Answer = {}  
for x1 in R1 do  
    for x2 in R2 do  
        ....  
        for xn in Rn do  
            if Conditions  
                then Answer = Answer  $\cup$  {(a1,...,ak)}  
return Answer
```

Meaning (Semantics) of SQL Queries

```
SELECT a1, a2, ..., ak  
FROM   R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE  Conditions
```

2. Parallel assignment

- Answer = {}
- **for** all assignments x1 **in** R1, ..., xn **in** Rn **do**
- **if** Conditions **then** Answer = Answer \cup {(a1,...,ak)}
- **return** Answer

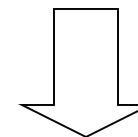
Doesn't impose any order!

Renaming Columns

•Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
•SELECT Pname AS prodName, Price AS askPrice  
FROM Product  
WHERE Price > 100
```



•Query with renaming

prodName	askPrice
SingleTouch	\$149.99
MultiTouch	\$203.99

Union, Intersection, Difference

```
(SELECT name  
FROM Person  
WHERE City="Seattle")
```

UNION

```
(SELECT name  
FROM Person, Purchase  
WHERE buyer=name AND store="The Bon")
```

- Similarly, you can use **INTERSECT** and **EXCEPT**.
- You must have the same attribute names (otherwise: rename).

```
•(SELECT DISTINCT R.A  
FROM R)  
INTERSECT ( (SELECT S.A FROM S)  
UNION  
(SELECT T.A FROM T))
```

Conserving Duplicates

```
(SELECT Name  
FROM Person  
WHERE City="Seattle")  
  
UNION ALL  
  
(SELECT Name  
FROM Person, Purchase  
WHERE Buyer=Name AND Store="The Bon")
```

Removing Duplicates

```
SELECT Company.Name  
FROM Company, Product, Purchase  
WHERE Company.Name= Product.Maker  
      AND Product.Name = Purchase.Product  
      AND Purchase.Buyer = 'Joe Blow'
```

- Multiple copies

```
SELECT DISTINCT Company.Name  
FROM Company, Product, Purchase  
WHERE Company.Name= Product.Maker  
      AND Product.Name = Purchase.product  
      AND Purchase.Buyer = 'Joe Blow'
```

- Single copies

Removing Duplicates

- **SELECT DISTINCT** Company.name
- **FROM** Company, Product
- **WHERE** Company.Name= Product.Maker
- AND Product.Name IN
- (**SELECT** Purchase.Product
- **FROM** Purchase
- **WHERE** Purchase.Buyer = ‘Joe Blow’)

- **SELECT DISTINCT** Company.Name
- **FROM** Company, Product, Purchase
- **WHERE** Company.Name= Product.Maker
- AND Product.Name = Purchase.Product
- AND Purchase.Buyer = ‘Joe Blow’

- Now
- they are
- equivalent

Aggregation

```
SELECT AVG(Price)  
FROM Product  
WHERE Maker="Toyota"
```

- SQL supports several aggregation operations:
 - SUM, MIN, MAX, AVG, COUNT

Aggregation: Count

```
•SELECT COUNT(*)  
•FROM Product  
•WHERE Year > 1995
```

- Except COUNT, all aggregations apply to a single attribute

Aggregation: Count

- COUNT applies to duplicates, unless otherwise stated:

- **SELECT** Count(Category) same as Count(*)
- **FROM** Product
- **WHERE** Year > 1995

- Better:

- **SELECT** Count(**DISTINCT** Category)
- **FROM** Product
- **WHERE** Year > 1995

Simple Aggregation

- Purchase(product, date, price, quantity)

- Example 1: find total sales for the entire database

```
SELECT SUM(price * quantity)  
FROM Purchase
```

- Example 1: find total sales of bagels

```
SELECT SUM(price * quantity)  
FROM Purchase  
WHERE product = 'bagel'
```

Grouping and Aggregation

- Usually, we want aggregations on certain parts of the relation.
- Purchase(product, date, price, quantity)
- Example 2: **find total sales after 10/1 per product.**

```
SELECT      product, SUM(price * quantity) AS TotalSales  
FROM        Purchase  
WHERE       date > "10/1"  
GROUPBY    product
```

Grouping and Aggregation

1. Compute the **FROM** and **WHERE** clauses.
 2. Group by the attributes in the **GROUPBY**
 3. Select one tuple for every group (and apply aggregation)
-
- SELECT** can have (1) grouped attributes or (2) aggregates.

Outer Joins

- Left outer join:
 - Include the left tuple even if there's no match
- Right outer join:
 - Include the right tuple even if there's no match
- Full outer join:
 - Include the both left and right tuples even if there's no match

Modifying the Database

Three kinds of modifications

- Insertions
- Deletions
- Updates

Sometimes they are all called “updates”

Insertions

- General form:

```
INSERT INTO R(A1,...,An) VALUES (v1,...,vn)
```

- Example: Insert a new purchase to the database:

```
INSERT INTO Purchase(buyer, seller, product, store)  
VALUES ('Joe', 'Fred', 'wakeup-clock-espresso-machine',  
       'The Sharper Image')
```

- Missing attribute → NULL.
- May drop attribute names if give them in order.

Insertions

```
INSERT INTO PRODUCT(name)
SELECT DISTINCT Purchase.product
FROM Purchase
WHERE Purchase.date > "10/26/01"
```

- The query replaces the VALUES keyword.
- Here we insert *many* tuples into PRODUCT

Insertion: an Example

- Product(name, listPrice, category)
- Purchase(prodName, buyerName, price)

- prodName is foreign key in Product.name

- Suppose database got corrupted and we need to fix it:

•Product

name	listPrice	category
gizmo	100	gadgets

•Purchase

prodName	buyerName	price
camera	John	200
gizmo	Smith	80
camera	Smith	225

- Task: insert in Product all prodNames from Purchase

Insertion: an Example

```
INSERT INTO Product(name)
SELECT DISTINCT prodName
FROM Purchase
WHERE prodName NOT IN (SELECT name FROM Product)
```

name	listPrice	category
gizmo	100	Gadgets
camera	-	-

Insertion: an Example

- **INSERT INTO** Product(name, listPrice)
- **SELECT DISTINCT** prodName, price
- **FROM** Purchase
- **WHERE** prodName **NOT IN** (**SELECT** name **FROM** Product)

name	listPrice	category
gizmo	100	Gadgets
camera	200	-
camera ??	225 ??	-

← Depends on the implementation

Deletions

- Example:

```
•DELETE FROM PURCHASE  
•WHERE seller = 'Joe' AND  
•          product = 'Brooklyn Bridge'
```

- Factoid about SQL: there is no way to delete only a single occurrence of a tuple that appears twice in a relation.

Updates

- Example:

```
UPDATE PRODUCT
SET price = price/2
WHERE Product.name IN
      (SELECT product
       FROM Purchase
       WHERE Date ='Oct, 25, 1999');
```

Data Definition in SQL

- So far we have seen the *Data Manipulation Language*, DML
- Next: *Data Definition Language* (DDL)
 - Data types:
 - Defines the types.
 - Data definition: defining the schema.
 - Create tables
 - Delete tables
 - Modify table schema
 - Indices: to improve performance

Data Types in SQL

- Characters:
 - CHAR(20) -- fixed length
 - VARCHAR(40) -- variable length
- Numbers:
 - INT, REAL plus variations
- Times and dates:
 - DATE, DATETIME (SQL Server only)
- To reuse domains:
CREATE DOMAIN address AS VARCHAR(55)

Creating Tables

- Example:

```
CREATE TABLE Person(  
    name          VARCHAR(30),  
    social-security-number INT,  
    age           SHORTINT,  
    city          VARCHAR(30),  
    gender        BIT(1),  
    Birthdate     DATE  
)
```

Deleting or Modifying a Table

- Deleting:

- Example: • **DROP Person;**
- Exercise with care !!

- Altering: (adding or removing an attribute).

- Example:

- **ALTER TABLE Person**
- **ADD phone CHAR(16);**
- **ALTER TABLE Person**
- **DROP age;**

- What happens when you make changes to the schema?

Default Values

- Specifying default values:

```
•CREATE TABLE Person(  
    •      name          VARCHAR(30),  
    •      social-security-number INT,  
    •      age           SHORTINT DEFAULT 100,  
    •      city          VARCHAR(30) DEFAULT 'Seattle',  
    •      gender        CHAR(1)   DEFAULT '?',  
    •      Birthdate     DATE
```

- The default of defaults: NULL

Defining Views

- Views are relations, except that they are not physically stored.
- For presenting different information to different users
- **Employee**(ssn, name, department, project, salary)
 - **CREATE VIEW** Developers **AS**
 - **SELECT** name, project
 - **FROM** Employee
 - **WHERE** department = “Development”
- Payroll has access to **Employee**, others only to **Developers**

A Different View

- Person(name, city)
- Purchase(buyer, seller, product, store)
- Product(name, maker, category)

```
• CREATE VIEW Seattle-view AS
  •   SELECT buyer, seller, product, store
  •   FROM Person, Purchase
  •   WHERE Person.city = "Seattle" AND
        Person.name = Purchase.buyer
```

- We have a new virtual table:
- Seattle-view(buyer, seller, product, store)

A Different View

- We can later use the view:

```
•SELECT name, store  
•FROM Seattle-view, Product  
•WHERE Seattle-view.product = Product.name AND  
•      Product.category = "shoes"
```

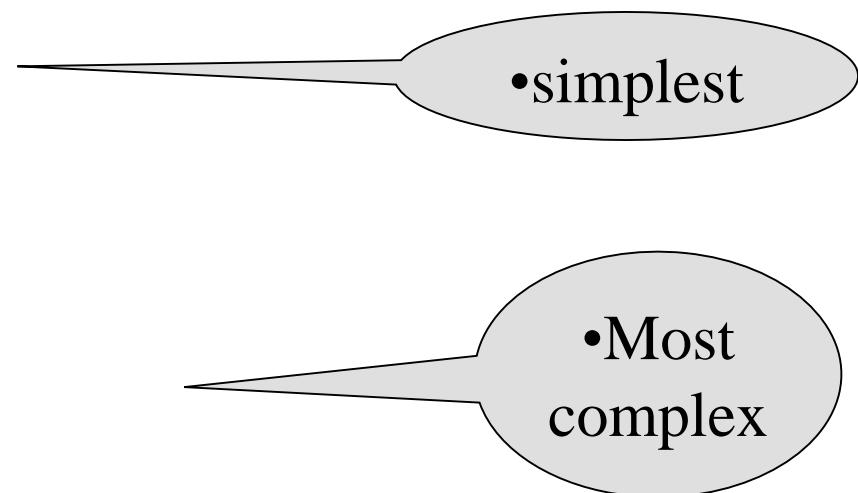
Constraints in SQL

- A constraint = a property that we'd like our database to hold
- The system will enforce the constraint by taking some actions:
 - forbid an update
 - or perform compensating updates

Constraints in SQL

Constraints in SQL:

- Keys, foreign keys
- Attribute-level constraints
- Tuple-level constraints
- Global constraints: assertions



The more complex the constraint, the harder it is to check and to enforce

Keys

```
•CREATE TABLE Product (
    •      name CHAR(30) PRIMARY KEY,
    •      category VARCHAR(20))
```

OR:

```
•CREATE TABLE Product (
    •      name CHAR(30),
    •      category VARCHAR(20)
•PRIMARY KEY (name))
```

Keys with Multiple Attributes

```
CREATE TABLE Product (
    name CHAR(30),
    category VARCHAR(20),
    price INT,
    PRIMARY KEY (name, category))
```

Other Keys

```
CREATE TABLE Product (
    productID CHAR(10),
    name CHAR(30),
    category VARCHAR(20),
    price INT,
    PRIMARY KEY (productID),
    UNIQUE (name, category))
```

- There is at most one **PRIMARY KEY**;
there can be many **UNIQUE**

Foreign Key Constraints

- Referential integrity constraints

```
CREATE TABLE Purchase (
    prodName CHAR(30)
        REFERENCES Product(name),
    date DATETIME)
```

- prodName is a **foreign key** to Product(name)
name must be a **key** in Product

•Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

•Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Foreign Key Constraints

- OR

```
•CREATE TABLE Purchase (
    • prodName CHAR(30),
    • category VARCHAR(20),
    • date DATETIME,
    • FOREIGN KEY (prodName, category)
    • REFERENCES Product(name, category)
```

- (name, category) must be a PRIMARY KEY

What happens during updates ?

Types of updates:

- In Purchase: insert/update
- In Product: delete/update

•Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

•Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Transactions

Address two issues:

- Access by multiple users
 - Remember the “client-server” architecture: one server with many clients
- Protection against crashes

Flight Reservation

- get values for :flight, :date, :seat
- EXEC SQL SELECT occupied INTO :occ
- FROM Flight
- WHERE fltNum = :flight AND fltdt= :date AND fltSeat=:seat
- **if** (!occ) {
- EXEC SQL UPDATE Flights
- SET occupied = ‘true’
- WHERE fltNum= :flight AND fltdt= :date AND fltSeat=:seat
- /* more code missing */
- }
- **else** /* notify customer that seat is not available */

Problem #1

- Customer 1 - finds a seat empty
- Customer 2 - finds the same seat empty
- Customer 1 - reserves the seat.
- Customer 2 - reserves the seat.
- Customer 1 will not be happy.
 - serializability*

Bank Transfers

- Transfer :amount from :account1 to :account2
- **EXEC SQL SELECT** balance INTO :balance1
 - FROM Accounts
 - WHERE accNo = :account1
- if (balance1 >= amount)
 - **EXEC SQL UPDATE** Accounts
 - SET balance = balance + :amount
 - WHERE acctNo = :account2;
 - **EXEC SQL UPDATE** Accounts
 - SET balance = balance - :amount
 - WHERE acctNo = :account1;

Transactions

- The user/programmer can group a sequence of commands so that they are executed atomically and in a serializable fashion:
 - **Transaction commit:** all the operations should be done and recorded.
 - **Transaction abort:** none of the operations should be done.
- In SQL:
 - EXEC SQL COMMIT;
 - EXEC SQL ROLLBACK;
- *Easier said than done...*

ACID Properties

- **Atomicity:** **all** actions of a transaction happen, or **none** happen.
- **Consistency:** if a transaction is consistent, and the database starts
 - from a consistent state, then it will end in a consistent
 - state.
- **Isolation:** the execution of one transaction is isolated from other
 - transactions.
- **Durability:** if a transaction commits, its effects persist in the
 - database.

How Do We Assure ACID?

- **Concurrency control:**

- Guarantees consistency and isolation, given atomicity.

- **Logging and Recovery:**

- Guarantees atomicity and durability.
*If you are going to be in the logging business, one of the things
that you'll have to do is learn about heavy equipment.*

-- Robert VanNatta

Logging History of Columbia County

Transactions in SQL

- In “ad-hoc” SQL:
 - Default: each statement = one transaction
- In “embedded” SQL:
 - BEGIN TRANSACTION
 - [SQL statements]
 - COMMIT or ROLLBACK (=ABORT)

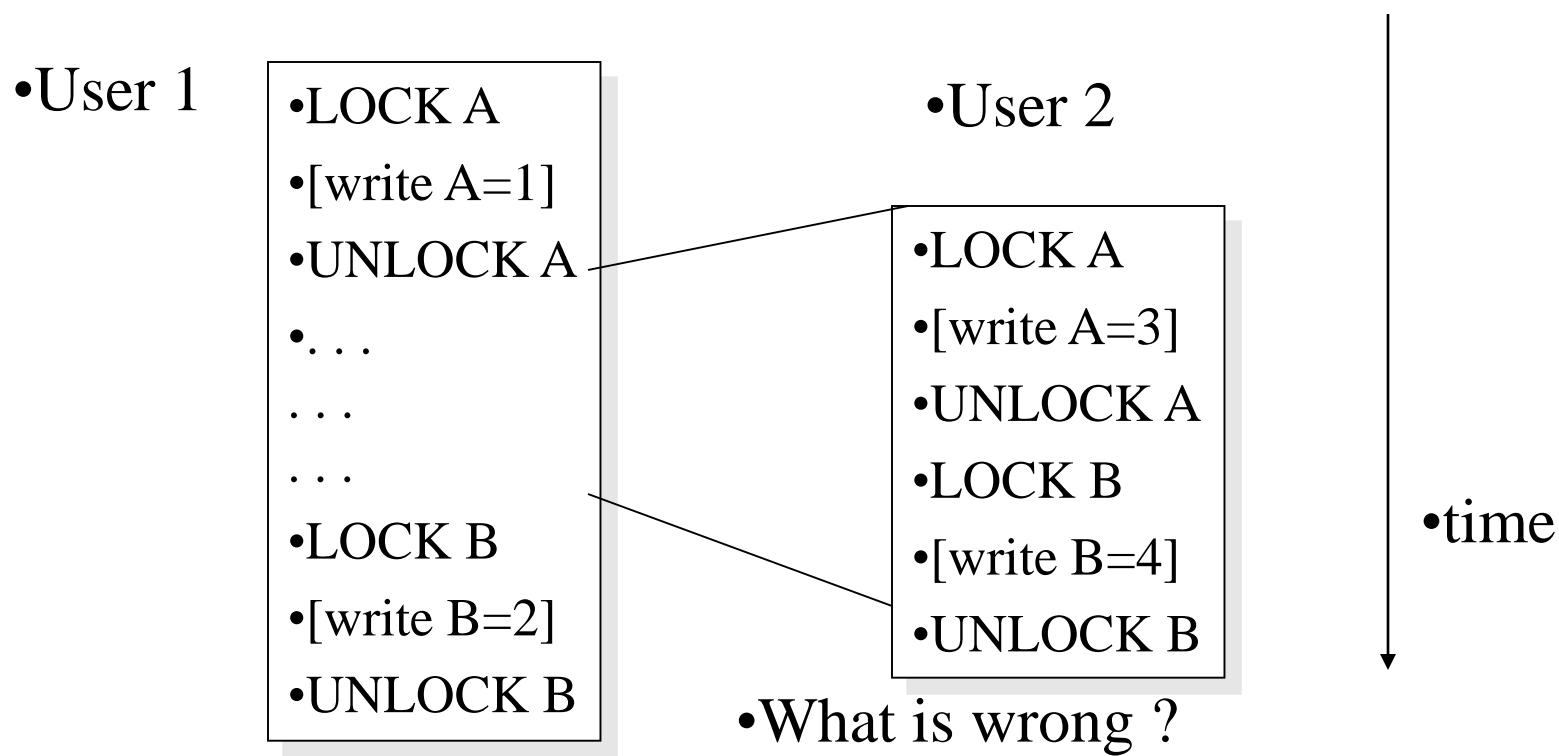
Transactions: Serializability

Serializability = the technical term for isolation

- An execution is ***serial*** if it is completely before or completely after any other transaction's execution
- An execution is ***serializable*** if it equivalent to one that is serial
- DBMS can offer serializability guarantees

Serializability

- Enforced with locks, like in Operating Systems!
- But this is not enough:



Serializability

- Solution: two-phase locking
 - Lock everything at the beginning
 - Unlock everything at the end
- Read locks: many simultaneous read locks allowed
- Write locks: only one write lock allowed
- Insert locks: one per table

Isolation Levels in SQL

1. “Dirty reads”

SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED

2. “Committed reads”

SET TRANSACTION ISOLATION LEVEL READ COMMITTED

3. “Repeatable reads”

SET TRANSACTION ISOLATION LEVEL REPEATABLE READ

4. Serializable transactions (default):

SET TRANSACTION ISOLATION LEVEL SERIALIZABLE

//TODO before next lecture:

- Homework 4 due on 4/16 at 11:59 pm EDT. Must be submitted on Submitty.