

Application Design Using Java

Lecture 23

Main Topics

- Overview of graph databases
- Installing and using Neo4j

Neo4j API

- HTTP API
 - Transactional Cypher HTTP endpoint
 - POST to a HTTP URL to send queries, and to receive responses from Neo4j
- Drivers
 - The preferred way to access a Neo4j server from an application
 - Use the Bolt protocol and have uniform design and use
 - Available in four languages: C# .NET, Java, JavaScript, and Python
 - Additional community drivers for: Spring, Ruby, PHP, R, Go, Erlang / Elixir, C/C++, Clojure, Perl, Haskell
 - API is defined independently of any programming language
- Procedures
 - Allow Neo4j to be extended by writing custom code which can be invoked directly from Cypher
 - Written in Java and compiled into jar files
 - To call a stored procedure, use a Cypher CALL clause

Neo4j Resources

- Neo4j Web site: <https://neo4j.com/>
- Neo4j installation manual: <https://neo4j.com/docs/operations-manual/current/deployment/single-instance/>
- Cypher Refcard <https://neo4j.com/docs/cypher-refcard/current/>
- Coursera course “Graph Analytics for Big Data” from the University of California, San Diego (<https://www.coursera.org/learn/big-data-graph-analytics>) has a lesson “Graph Analytics With Neo4j”
- Webber, Jim. "A programmatic introduction to Neo4j." *Proceedings of the 3rd annual conference on Systems, programming, and applications: software for humanity*. ACM, 2012.
- Robinson, Ian, James Webber, and Emil Eifrem. *Graph databases*. Sebastopol, CA: O'Reilly, 2015
- Bruggen, Rik. *Learning Neo4j*. Birmingham, UK: Packt Pub, 2014

Main Topics

- Overview of graph databases
- Installing and using Neo4j

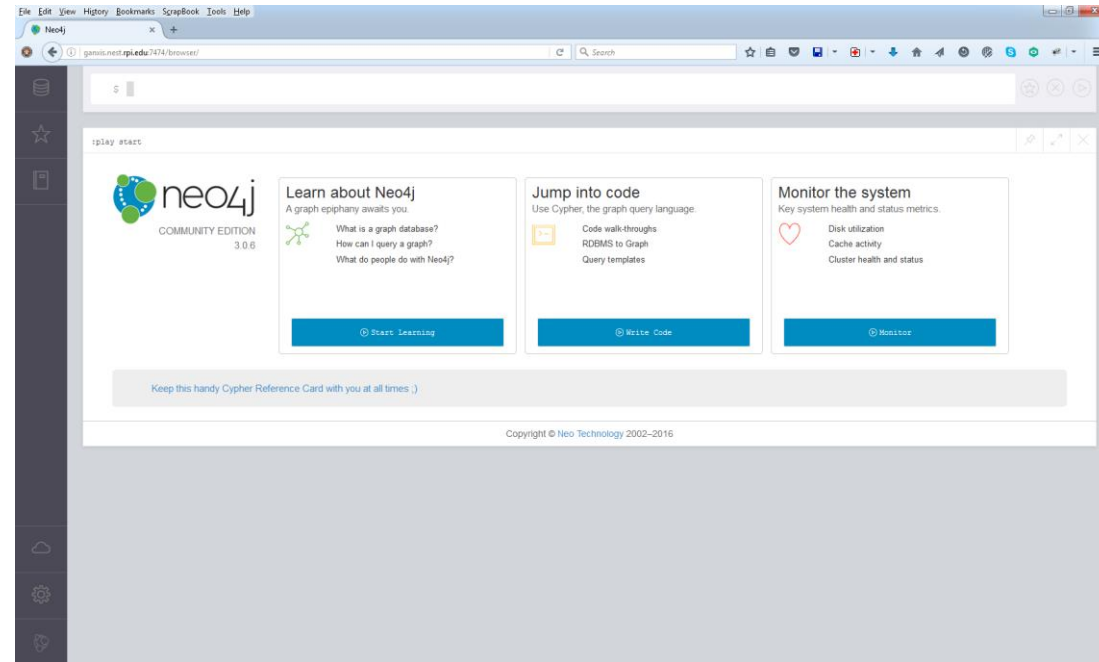
Neo4j Installation

- Neo4j runs on Linux, Windows, and OS X
- A Java 8 runtime is required
- Several ways to install on Linux, depending on the Linux distro (see the “Neo4j Resources” slide)
- Check the `/etc/neo4j/neo4j.conf` configuration file:

```
# HTTP Connector
dbms.connector.http.type=HTTP
dbms.connector.http.enabled=true
# To accept non-local HTTP connections, uncomment this line
dbms.connector.http.address=0.0.0.0:7474
```
- File locations depend on the operating system, as described here:
<https://neo4j.com/docs/operations-manual/current/deployment/file-locations/>
- Make sure you start the Neo4j server (e.g., “`./bin/neo4j start`” or “`service neo4j start`” on Linux)

Neo4j Browser

- Open the URL <http://localhost:7474> (replace “localhost” with your server name, and 7474 with the port name as set in neo4j.conf)
- Enter the username/password (if not set, Neo4j browser will prompt you to select the username and password)
- Start working with Neo4j by entering Cypher queries and observing their results
- Save frequently used Queries to Favorites



The Structure of a Cypher Query

- Nodes are surrounded with parentheses which look like circles, e.g. (a)
- A relationship is basically an arrow --> between two nodes with additional information placed in square brackets inside of the arrow
- A query is comprised of several distinct clauses, like:
 - MATCH: The graph pattern to match. This is the most common way to get data from the graph.
 - WHERE: Not a clause in its own right, but rather part of MATCH, OPTIONAL MATCH and WITH. Adds constraints to a pattern, or filters the intermediate result passing through WITH.
 - RETURN: What to return.

Cypher using relationship 'likes'



Cypher

(a) -[:LIKES]-> (b)

```
MATCH (john {name: 'John'})-[:friend]->()-[:friend]->(fof) RETURN john.name, fof.name
```



Writing Cypher Queries

- Node labels, relationship types and property names are **case-sensitive** in Cypher
- **CREATE** creates nodes with labels and properties or more complex structures
- **MERGE** matches existing or creates new nodes and patterns. This is especially useful together with uniqueness constraints.
- **DELETE** deletes nodes, relationships, or paths. Nodes can only be deleted when they have no other relationships still existing
- **DETACH DELETE** deletes nodes and all their relationships
- **SET** sets values to properties and add labels on nodes
- **REMOVE** removes properties and labels on nodes
- **ORDER BY** is a sub-clause that specifies that the output should be sorted and how

Importing and Exporting Data

- Loading data from CSV is the most straightforward way of importing data into Neo4j
- For fast batch import of huge datasets, use the neo4j-import tool
- Lots of other tools for different data formats and database sizes
- More on importing data at <https://neo4j.com/developer/guide-importing-data-and-etl/>
- Export data using Neo4j browser or neo4j-shell-tools

Loading Data from CSV



- Understand your graph model

```
(p1:Person {userId:10,name:"Anne"})-[:KNOWS]->(p2:Person {userId:123,name:"John"})
```

- CSV files

- people.csv
1,"John" 10,"Jane" 234,"Fred" 4893,"Mark" 234943,"Anne"
- friendships.csv
1,234 10,4893 234,1 4893,234943 234943,234 234943,1

- Run the following Cypher queries:

- CREATE CONSTRAINT ON (p:Person) ASSERT p.userId IS UNIQUE;
- LOAD CSV FROM "file:///people.csv" AS csvLine
MERGE (p:Person {userId: toInteger(csvLine[0])}, name: csvLine[1]);
- LOAD CSV FROM "file:///friendships.csv" AS csvLine MATCH (p1:Person {userId: toInteger(csvLine [0])}), (p2:Person {userId: toInteger(csvLine [1])})
CREATE (p1)-[:KNOWS]->(p2);
- CREATE INDEX ON :Person(name);

- Check the results:

```
MATCH (:Person {name:"Anne"})-[:KNOWS*2..2]-(p2) RETURN p2.name, count(*) as freq  
ORDER BY freq DESC;
```

Loading Data from a Spreadsheet

- Lay out your data in a spreadsheet
- Use formulas to generate the required Cypher statements

	A	B	C	D	E
	Browser				
	id	name	gender	from-name	to-name
1	Meetup Malmö				
2	Donald Duck	man		2	5
3	Fredrik	man		7	1
4	Jakub	man		5	10
5	Johan	man		8	9
6	Patrik	man		3	4
7	Pernilla	woman		2	12
8	Peter	man		3	2
9	Rik	man		9	8
10	Sebastian	man		7	3
11	Thorbjörn	man		11	8
12	Tobias	man		12	8

fx	=merge (_"&A4&":Person {id:'"&A4&"', name:'"&B4&"', gender:'"&C4&"'})-[:ATTENDS]->(meetup)"									
	A	B	C	D			E	F	G	H
1	Browser									
2	id	name	gender				from-	to-nan	type	Relationship
3	1	Meetup Malmö		merge (meetup:Event {id:'1', name:'Meetup Malmö'})						
4	2	Donald Duck	man	merge (2:Person {id:'2', name:'Donald Duck', gender:'man'})-[:ATTENDS]->(meetup)			2	5	WORKS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'2'}), (p2:Person {id:'5'}) MERGE (p1)-[V
5	3	Fredrik	man	merge (3:Person {id:'3', name:'Fredrik', gender:'man'})-[:ATTENDS]->(meetup)			7	1	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'7'}), (p2:Person {id:'1'}) MERGE (p1)-[F
6	4	Jakub	man	merge (4:Person {id:'4', name:'Jakub', gender:'man'})-[:ATTENDS]->(meetup)			5	10	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'5'}), (p2:Person {id:'10'}) MERGE (p1)-[F
7	5	Johan	man	merge (5:Person {id:'5', name:'Johan', gender:'man'})-[:ATTENDS]->(meetup)			8	9	WORKS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'8'}), (p2:Person {id:'9'}) MERGE (p1)-[V
8	6	Patrik	man	merge (6:Person {id:'6', name:'Patrik', gender:'man'})-[:ATTENDS]->(meetup)			3	4	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'3'}), (p2:Person {id:'4'}) MERGE (p1)-[F
9	7	Pernilla	woman	merge (7:Person {id:'7', name:'Pernilla', gender:'woman'})-[:ATTENDS]->(meetup)			2	12	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'2'}), (p2:Person {id:'12'}) MERGE (p1)-[F
10	8	Peter	man	merge (8:Person {id:'8', name:'Peter', gender:'man'})-[:ATTENDS]->(meetup)			3	2	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'3'}), (p2:Person {id:'2'}) MERGE (p1)-[F
11	9	Rik	man	merge (9:Person {id:'9', name:'Rik', gender:'man'})-[:ATTENDS]->(meetup)			9	8	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'9'}), (p2:Person {id:'8'}) MERGE (p1)-[F
12	10	Sebastian	man	merge (10:Person {id:'10', name:'Sebastian', gender:'man'})-[:ATTENDS]->(meetup)			7	3	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'7'}), (p2:Person {id:'3'}) MERGE (p1)-[F
13	11	Thorbjörn	man	merge (11:Person {id:'11', name:'Thorbjörn', gender:'man'})-[:ATTENDS]->(meetup)			11	8	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'11'}), (p2:Person {id:'8'}) MERGE (p1)-[F
14	12	Tobias	man	merge (12:Person {id:'12', name:'Tobias', gender:'man'})-[:ATTENDS]->(meetup)			12	8	FRIENDS_WITH	WITH 1 as dummy MATCH (p1:Person {id:'12'}), (p2:Person {id:'8'}) MERGE (p1)-[F
15										

- Collect Cypher queries and run them
- Check the results:

```
MATCH (p1:Person) -[:ATTENDS]- (e:Event {name:"Meetup Malmö"}) -[:ATTENDS]-
(p2:Person) WHERE (p1)-[:FRIENDS_WITH]- (p2) RETURN p1, p2, e;
```

Loading Data from a GraphML file

- Use **neo4j-shell-tools** from <https://github.com/jexp/neo4j-shell-tools>
- Populate the database from a GraphML file

```
import-graphml -i /usr/share/neo4j/import/airlines.graphml -r  
HAS_DIRECT_FLIGHTS_TO -b 20000 -c -t
```

- Check the results:

```
MATCH (a)--()  
WITH a.tooltip as airport, count(*) as flights  
RETURN airport, flights ORDER BY flights DESC LIMIT 10
```

Loading Data from an Arbitrary Format

- Write a simple program to convert your file into a set of two CSV files
- Load data from the CSV file into a Neo4j database
 - `CREATE CONSTRAINT CharacterNameUnique ON (c:Character) ASSERT c.name IS UNIQUE;`
 - `LOAD CSV WITH HEADERS FROM "file:///Marvel-nodes.csv" AS csvLine MERGE (c:Character {name: csvLine.NodeID});`
 - `LOAD CSV WITH HEADERS FROM "file:///Marvel-edges.csv" AS csvLine MATCH (c1:Character {name: csvLine.EdgeFrom}), (c2:Character {name: csvLine.EdgeTo}) CREATE (c1)-[:APPEARED_WITH]->(c2);`
 - `CREATE INDEX FOR (c:Character) ON (c.name);`
- Check the results:


```
MATCH (c:Character)-[r]-()
WITH c as characters, count(distinct r) as degree
RETURN degree, count(characters) ORDER BY degree ASC
```

```
from sys import argv

def read_edge_list(filename):
    nodeset= set([])
    edgelist = []
    with open(filename, 'r') as file_handle:
        for line in file_handle:
            if line[0] != '#':
                data = line.split(',')
                node_from = data[0] + ''
                node_to = '' + data[1].strip()
                nodeset.add(node_from)
                nodeset.add(node_to)
                edgelist.append([node_from, node_to])
    return nodeset, edgelist


def write_csv_nodes(nodes, file_nodes):
    with open(file_nodes, 'w') as file_handle:
        file_handle.write("NodeID\n")
        for node in nodes:
            file_handle.write('{0}\n'.format(node))

def write_csv_edges(edges, file_nodes):
    with open(file_nodes, 'w') as file_handle:
        file_handle.write("EdgeFrom,EdgeTo\n")
        for edge in edges:
            file_handle.write('{0},{1}\n'.format(edge[0], edge[1]))

script, input_file, output_file_nodes,
output_file_edges = argv
nodes, edges = read_edge_list(input_file)
write_csv_nodes(nodes, output_file_nodes)
write_csv_edges(edges, output_file_edges)
```

Exporting Data From Neo4j

- Click the download icon on the table view of the Cypher query results
- Use neo4j-shell-tools to export results of a Cypher query to a CSV or GraphML file
- Access the graph data with Neo4j API and save it in the desired format






Analyzing Graph Data with MATLAB

- Load CSV data exported from Neo4j into MATLAB
- Use MATLAB to perform additional analysis and to draw plots
- Export analysis results and plots for publication

```
filename = 'filename.csv';  
M = csvread(filename,1,0);  
x = M(:,1);  
y = M(:,2);  
plot(x,y)  
f=fit(x,y,'poly2')  
plot(f,x,y)  
f=fit(x,y,'power1')  
plot(f,x,y)
```

```
f =  
  
General model Power1:  
f(x) = a*x^b  
Coefficients (with 95% confidence bounds):  
a =      2282   (2257, 2308)  
b =     -1.579   (-1.607, -1.552)
```



Using Transactional Cypher HTTP Endpoint

- Allows you to execute a series of Cypher statements within the scope of a transaction
- The transaction may be kept open across multiple HTTP requests, until the client chooses to commit or roll back
- Each HTTP request can include a list of statements
- Requests should include an Authorization header, with a value of Basic <payload>, where "payload" is a base64 encoded string of "username:password"

```
import requests
from requests.exceptions import ConnectionError
import json

NEO4J_SERVER = 'http://localhost:7474'
NEO4J_COMMIT_ENDPOINT = '/db/data/transaction/commit'
NEO4J_CREDENTIALS = '*****'

def execute_neo4j_cypher(url, credentials, query, parameters):
    result = None
    query_text = json.dumps(dict(statements = [dict(statement =
query, parameters = parameters)]))

    headers = {'Accept' : 'application/json', 'Content-type' :
'application/json', 'Authorization:' : 'Basic ' + credentials}
    try:
        resp = requests.post(url, headers = headers, data =
query_text)
        result = resp.json()
    except ConnectionError as exception:
        print exception # Log error
    if len(result['errors']) > 0:
        print '*** ERROR! Error executing Cypher query' # Log
error
        print '*** ', query, '<-', parameters
        print '*** ' + str(result)
    return result

query = 'MERGE (p: Person {id:{userid}, name:{name}}) ON CREATE
SET p.created = timestamp() ON MATCH SET p.matched =
timestamp() RETURN p'
parameters = dict()
parameters['userid'] = 17
parameters['name'] = 'J J'
execute_neo4j_cypher(NEO4J_SERVER + NEO4J_COMMIT_ENDPOINT,
NEO4J_CREDENTIALS, query, parameters)
```

Using Drivers to Access Neo4j

- Binary Bolt protocol (starting with Neo4j 3.0)
- Binary protocol is enabled in Neo4j by default and can be used in any language driver that supports it
- Native Java driver officially supported by Neo4j
- Drivers implement all low level connection and communication tasks

```
import org.neo4j.driver.v1.*;

public class Neo4j
{
    public static void javaDriverDemo() {
        Driver driver = GraphDatabase.driver("bolt://localhost", "neo4j", "neo4j");
        Session session = driver.session();

        StatementResult result = session.run("MATCH (a)-[]-(b)-[]-(c)-[]-(a) WHERE a.id < b.id AND b.id < c.id RETURN DISTINCT a,b,c");
        int counter = 0;
        while (result.hasNext())
        {
            counter++;
            Record record = result.next();
            System.out.println(record.get("a").get("id") + " \t" + record.get("b").get("id") + " \t" + record.get("c").get("id"));
        }
        System.out.println("Count: " + counter);
        session.close();
        driver.close();
    }

    public static void main(String [] args)
    {
        javaDriverDemo();
    }
}
```

Using Core Java API

- Native Java API performs database operations directly with Neo4j core

```
import java.io.*;
import java.util.*;
import org.neo4j.graphdb.*

public class Neo4j
{
    public enum NodeLabels implements Label { NODE; }
    public enum EdgeLabels implements RelationshipType{ CONNECTED; }
    public static void javaNativeDemo(int nodes, double p) {
        Node node1, node2; Random randomgen = new Random();
        GraphDatabaseFactory dbFactory = new GraphDatabaseFactory();
        GraphDatabaseService db = dbFactory.newEmbeddedDatabase(new File("TestNeo4jDB"));
        try (Transaction tx = db.beginTx()) {
            for (int i = 1; i <= nodes; i++) {
                Node node = db.createNode(NodeLabels.NODE);
                node.setProperty("id", i);
            }
            for (int i = 1; i <= nodes; i++)
                for (int j = i + 1; j <= nodes; j++) {
                    if (randomgen.nextDouble() < p) {
                        node1 = db.findNode(NodeLabels.NODE, "id", i);
                        node2 = db.findNode(NodeLabels.NODE, "id", j);
                        Relationship relationship = node1.createRelationshipTo(node2, EdgeLabels.CONNECTED);
                        relationship = node2.createRelationshipTo(node1, EdgeLabels.CONNECTED);
                    }
                }
            tx.success();
        }
        db.shutdown();
    }
    public static void main(String [] args) {
        javaNativeDemo(100, 0.2);
    }
}
```

//TODO before next lecture:

- Homework 5 due on 4/30 at 11:59 pm EDT. Must be submitted on Submittity.